6502509794
chris.j.cundy@gmail.com
C-J-Cundy
Updated November 6, 2026

Chris Cundy

Experience

June 2024– **Research Scientist**, *FAR.AI*, Berkeley, California, USA

Characterising and mitigating catastrophic risks from frontier AI systems. My main responsibilities are:

- Leading research projects and determining the direction for teams of researchers and engineers.
- Designing and implementing training and evaluation schemes with state-of-the-art models, conducting empirical studies, and communicating results through papers and presentations.
- As a research team lead, ensuring that research aims are consistent with FAR's mission and that my direct reports are supported in their output and professional development.

June 2022–September 2022 **Research Scientist Intern**, *Technical AI Safety Team*, DeepMind, London, UK

Investigating robust and reliable machine learning in theory and at scale

- Investigated susceptibility of autoregressive models to 'delusions', where unobserved latent variables lead to incorrect probabilistic judgments.
- Developed a theoretical model for delusions; investigated delusions at scale by analysing performance of DeepMind's Gato (a large generalist, multi-task autoregressive model) on custom environments.

October 2017–January 2018 **Visiting Scholar**, *Future of Humanity Institute*, University of Oxford, Oxford, UK

Developing algorithms to predict human judgments. Supervised by Owain Evans and Andreas Stuhlmüller

- Designed algorithms to collate quick, noisy human judgments to predict the answer to complicated tasks which would typically require deliberation.

June–September 2017 **Visiting Scholar**, *Centre for Human-Compatible AI*, University of California, Berkeley, US

Supervised by Daniel Filan & Stuart Russell. I analysed inverse reinforcement learning in the hierarchical options setting. I formalized the problem, performed experiments on data and presented at an ICML workshop.

Education

2018-2024 **PhD - Computer Science**, *Stanford University*, Stanford, California, USA

- Advised by Stefano Ermon.
- Investigating topics in inverse reinforcement learning, sequence modelling and variational inference.
- Thesis: *Beyond Maximum Likelihood: Distribution-Aware Machine Learning*.

2016-2017 **MEng - Computer Science**, *University of Cambridge*, Cambridge, UK

- Grade: Distinction. Modules include: Data Science, Probabilistic Machine Learning, Network Analytics.
- Supervised by Carl E. Rasmussen.

2013-2016 **BA - Natural Sciences (Physics)**, *University of Cambridge*, Cambridge, UK

- Grade: 1st. Modules: Physics, Maths, Chemistry, Computer Science.

Selected Publications

2026 **Telling FIBS: Probes Effectively Detect Sabotage and Catch Unverbalized Deception**, *Oskar J. Hollinsworth, Alex F. Spies, Tigist Diriba, Adam Gleave,* **Chris Cundy**, Preprint

2026 **The Obfuscation Atlas: Mapping Where Honesty Emerges in RLVR with Deception Probes**,
Mohammad Taufeeque, Stefan Heimersheim, Adam Gleave, **Chris Cundy**, ICML 2026 (Outstanding paper honorable mention, top seven (0.1%) among accepted papers)

2025 **Auditing Games for Sandbagging**,
Jordan Taylor, Sid Black, Dillon Bowen, Thomas Read, Satvik Golechha, Alex Zelenka-Martin, Oliver Makins, Connor Kissane, Kola Ayonrinde, Jacob Merizian, Samuel Marks, **Chris Cundy**, *Joseph Bloom*, Technical Report

2025 **Preference Learning with Lie Detectors can Induce Honesty or Evasion**,
Chris Cundy, *Adam Gleave*, NeurIPS 2025

2024 **SequenceMatch: Imitation Learning for Autoregressive Sequence Modelling with Backtracking**,
Chris Cundy, *Stefano Ermon*, ICLR 2024

2024 **Privacy-Constrained Policies via Mutual Information Regularized Policy Gradients**,
Chris Cundy, *Rishi Desai, Stefano Ermon*, AISTATS 2024

2021 **BCD Nets: Scalable Variational Approaches for Bayesian Causal Discovery**,
Chris Cundy, *Aditya Grover, Stefano Ermon*, NeurIPS 2021

2020 **Flexible Approximate Inference via Stratified Normalizing Flows**,
Chris Cundy, *Stefano Ermon*, UAI 2020

2018 **Parallelizing Linear Recurrent Neural Nets over Sequence Length**,
Eric Martin, **Chris Cundy**, ICLR 2018

Additional Publications

2026 **Scaling Trends for Lie Detector Oversight in Preference Learning**,
Oskar J. Hollinsworth, Ann-Kathrin Dombrowski, Sam Adam-Day, Adam Gleave, **Chris Cundy**, AI4Good Workshop, ICML 2026

2026 **The Safety Gap Toolkit: Evaluating Hidden Dangers of Open-Source Models**,
Ann-Kathrin Dombrowski, Dillon Bowen, Adam Gleave, **Chris Cundy**, LockLLM Workshop, NeurIPS 2025

2025 **Sharpe Ratio-Guided Active Learning for Preference Optimization in RLHF**,
Syrine Belakaria, Joshua Kazdan, Charles Marx, **Chris Cundy**, *Willie Neiswanger, Sanmi Koyejo, Barbara E Engelhardt, Stefano Ermon*, CoLM 2025

2024 **A physics-informed machine learning model for the prediction of drop breakup in two-phase flows**,
Chris Cundy, *Shahab Mirjalili, Charlélie Laurent, Stefano Ermon, Gianluca Iaccarino, Ali Mani*, International Journal of Multiphase Flow

2023 **Neural Networks and the Chomsky Hierarchy**,
Grégoire Delétang, Anian Ruoss, Jordi Grau-Moya, Tim Genewein, Li Kevin Wenliang, Elliot Catt, **Chris Cundy**, *Marcus Hutter, Shane Legg, Joel Veness, Pedro A. Ortega*, ICLR 2023

2022 **Towards a foundation model for geospatial artificial intelligence**,
Gengchen Mai, **Chris Cundy**, *Kristy Choi, Yingjie Hu, Ni Lao, Stefano Ermon*, Proceedings of the 30th International Conference on Advances in Geographic Information Systems

2022 **LMPriors: Pre-Trained Language Models as Task-Specific Priors**,
*Kristy Choi**, **Chris Cundy***, *Sanjari Srivasta, Stefano Ermon*, First Workshop on Foundation Models for Decision Making, NeurIPS 2022

2021 **IQ-Learn: Inverse soft-Q Learning for Imitation**,
Divyansh Garg, Shuvam Chakraborty, **Chris Cundy**, *Jiaming Song, Stefano Ermon*, NeurIPS 2021

2018 **Exploring Hierarchy-Aware Inverse Reinforcement Learning**,
Chris Cundy, *Daniel Filan*, Workshop on Goal Specifications for Reinforcement Learning, ICML 2018

2017 **Predicting Slow Judgment**,
Owain Evans, Andreas Stuhlmüller, Ryan Carey, Neal Jean, Andrew Schreiber, Girish Sastry, **Chris Cundy**, Aligned Artificial Intelligence Workshop, NeurIPS 2017

Service

2025 **Participant, EU AI Act Code of Practice Working Groups 2 and 4**
Participated, as an independent expert, in working groups 2 and 4 for the development of the EU AI Act Code of Practice (CoP). I advocated, via oral presentation, for both pre- and post-mitigation model evaluations

2023 **Teaching Assistant–CS228 (Probabilistic Graphical Models)**, *Stanford University*

2022 **Head Teaching Assistant–CS228 (Probabilistic Graphical Models)**, *Stanford University*
Received award for excellence (awarded to top 5% of Teaching Assistants).

2023 **Project Supervisor**, *Supervised Project for Alignment Research (SPAR)*, Stanford AI Alignment
Supervised five undergraduates on a project finding scaling laws in prompt injections.

2021 **Project Supervisor**, *Undergraduate Research Program*, Stanford Existential Risk Initiative
Served as supervisor for an undergraduate project on forecasting AI progress.

2020– **Reviewer**
Reviewed for the following venues: UAI (2020,2022,2025,2026), ICML (2019,2020,2023,2025,2026), ICLR (2021-2026), NeurIPS (2021-2025), AAAI-(Safe and Robust AI track) (2023-2024), TMLR (2025,2026)

Relevant Awards

March 2024 **Winner, OpenAI Preparedness Challenge**

- One of the top ten submissions for the OpenAI Preparedness Challenge, for submitting *the most unique, while still being probable, potentially catastrophic misuse of the [OpenAI API]*.
- Developed proof-of-concept showing how GPT4-V, and speech-to-text with GPT4, could be used to parse vast amounts of unlabelled surveillance data, finding actionable insights for blackmail or insider trading.